

Time Series Fundamentals

Kaiwen Zhou

Contents

1	Topic: Mean Squared Prediction	1
2	Topic: MLE for AR and MA processes	2
3	Topic: MLE for AR and MA processes	3
4	Application: Mean Squared Prediction and the Best Linear Predictor	4

1 Topic: Mean Squared Prediction

Suppose $\{X_t\}$ is a stationary process with mean μ , and auto-covariance function $\gamma_X(h)$. Our goal is to predict: X_{n+h} using observations X_1, X_2, \dots, X_n .

We define the **Best Linear Predictor** to be the one minimizes the Mean Squared Prediction Error (MSPE):

$$\hat{X}_{n+h} = E[X_{n+h} | X_1, \dots, X_n] = P_n X_{n+h} := a_0^* + a_1^* X_n + \dots + a_n^* X_1 = a_0 + \mathbf{a}_n^\top \mathbf{x}_n = \arg \min_{a_0, \mathbf{a}_n} \mathbb{E} \left[(X_{n+h} - a_0 - \mathbf{a}_n^\top \mathbf{x}_n)^2 \right] = \arg \min_{a_0, \mathbf{a}_n} L(a_0, \mathbf{a}_n)$$

where $\mathbf{a}_n = (a_1, \dots, a_n) \in \mathbb{R}^n$, $\mathbf{x}_n = (x_n, \dots, x_1) \in \mathbb{R}^n$, P in P_n stands for “Predictor” and n suggests that we are using n most recent observations.

To simplify our computation, we define $\mathbf{1} = (1, \dots, 1) \in \mathbb{R}^n$, $\gamma_n(h) = \begin{bmatrix} \gamma_X(h) \\ \gamma_X(h+1) \\ \vdots \\ \gamma_X(h+n-1) \end{bmatrix} \in \mathbb{R}^n$, $\Gamma_n = \begin{bmatrix} \gamma_X(0) & \gamma_X(1) & \cdots & \gamma_X(n-1) \\ \gamma_X(1) & \gamma_X(0) & \cdots & \gamma_X(n-2) \\ \vdots & \vdots & \ddots & \vdots \\ \gamma_X(n-1) & \gamma_X(n-2) & \cdots & \gamma_X(0) \end{bmatrix} \in \mathbb{R}^{n \times n}$.

Taking the partial derivatives of L , we obtain

$$0 = \frac{\partial L}{\partial a_0} = \mathbb{E} \left[\frac{\partial}{\partial a_0} (X_{n+h} - a_0 - \mathbf{a}_n^\top \mathbf{x}_n)^2 \right] = \mathbb{E} \left[-2(X_{n+h} - a_0 - \mathbf{a}_n^\top \mathbf{x}_n) \right] = 2(a_0 - \mu + \mathbf{a}_n^\top \mathbf{1}) \implies a_0^* = \mu(1 - \mathbf{a}_n^\top \mathbf{1})$$

and

$$\begin{aligned} \mathbf{0} &= \frac{\partial L}{\partial \mathbf{a}_n} = \frac{\partial}{\partial \mathbf{a}_n} \mathbb{E} \left[(X_{n+h} - a_0 - \mathbf{a}_n^\top \mathbf{x}_n)^2 \right] \\ &= \mathbb{E} \left[\frac{\partial}{\partial \mathbf{a}_n} \left(\mathbf{a}_n^\top \mathbf{x}_n \mathbf{x}_n^\top \mathbf{a}_n - 2X_{n+h} \mathbf{x}_n^\top \mathbf{a}_n + 2a_0 \mathbf{x}_n^\top \mathbf{a}_n + (\text{terms without } \mathbf{a}_n) \right) \right] \\ &= 2\mathbb{E}[\mathbf{x}_n \mathbf{x}_n^\top \mathbf{a}_n - X_{n+h} \mathbf{x}_n + a_0 \mathbf{x}_n] \\ &= 2\mathbb{E}[\mathbf{x}_n \mathbf{x}_n^\top] \mathbf{a}_n - 2\mathbb{E}[X_{n+h} \mathbf{x}_n] + 2\mathbb{E}[a_0 \mathbf{x}_n] \\ &= 2(\Gamma_n + \mu^2 \mathbf{1} \mathbf{1}^\top) \mathbf{a}_n - 2(\gamma_n(h) + \mu^2 \mathbf{1}) + 2a_0 \mu \mathbf{1} \\ a_0 &= \mu(1 - \mathbf{a}_n^\top \mathbf{1}) \implies = 2(\Gamma_n + \mu^2 \mathbf{1} \mathbf{1}^\top) \mathbf{a}_n - 2(\gamma_n(h) + \mu^2 \mathbf{1}) + 2\mu(1 - \mathbf{a}_n^\top \mathbf{1}) \mu \mathbf{1} \\ &= 2\Gamma_n \mathbf{a}_n - 2\gamma_n(h) \\ &\implies \mathbf{a}_n^* = \Gamma_n^{-1} \gamma_n(h) \end{aligned}$$

Therefore, the **Best Linear Predictor** is given by

$$P_n X_{n+h} := a_0^* + \mathbf{a}_n^{*\top} \mathbf{x}_n \text{ where } \begin{cases} a_0^* &= \mu(1 - \mathbf{a}_n^{*\top} \mathbf{1}) \\ \mathbf{a}_n^* &= \Gamma_n^{-1} \gamma_n(h) \end{cases} \quad (1)$$

Problem 1.1. Show that Mean Squared Predictive Error is obtained by

$$\mathbb{E} \left[(X_{n+h} - P_n X_{n+h})^2 \right] = \gamma_X(0) - \mathbf{a}_n^{*\top} \gamma_n(h)$$

Solution: From above, we get

$$\begin{aligned} \mathbb{E} \left[(X_{n+h} - P_n X_{n+h})^2 \right] &= \mathbb{E} \left[(X_{n+h} - a_0^* - \mathbf{a}_n^{*\top} \mathbf{x}_n)^2 \right] \\ \mathbf{a}_n^{*\top} \mathbf{x}_n &= \mathbf{x}_n^\top \mathbf{a}_n^* \implies = \mathbb{E} \left[X_{n+h}^2 + a_0^{*2} + \mathbf{a}_n^{*\top} \mathbf{x}_n \mathbf{x}_n^\top \mathbf{a}_n^* - 2a_0^* X_{n+h} - 2X_{n+h} \mathbf{a}_n^{*\top} \mathbf{x}_n + 2a_0^* \mathbf{a}_n^{*\top} \mathbf{x}_n \right] \\ &= \mathbb{E} \left[X_{n+h}^2 \right] + a_0^{*2} + \mathbf{a}_n^{*\top} \mathbb{E}[\mathbf{x}_n \mathbf{x}_n^\top] \mathbf{a}_n^* - 2a_0^* \mathbb{E}[X_{n+h}] - 2\mathbf{a}_n^{*\top} \mathbb{E}[X_{n+h} \mathbf{x}_n] + 2a_0^* \mathbf{a}_n^{*\top} \mathbb{E}[\mathbf{x}_n] \\ &= \gamma_X(0) + \mu^2 + \mu^2(1 - \mathbf{a}_n^{*\top} \mathbf{1})^2 + \mathbf{a}_n^{*\top} (\Gamma_n + \mu^2 \mathbf{1} \mathbf{1}^\top) \mathbf{a}_n^* - 2\mu^2(1 - \mathbf{a}_n^{*\top} \mathbf{1}) - 2\mathbf{a}_n^{*\top} (\gamma_n(h) + \mu^2 \mathbf{1}) + 2\mu^2(1 - \mathbf{a}_n^{*\top} \mathbf{1}) \mathbf{a}_n^{*\top} \mathbf{1} \\ &= \gamma_X(0) + \mathbf{a}_n^{*\top} \Gamma_n \mathbf{a}_n^* - 2\mathbf{a}_n^{*\top} \gamma_n(h) \\ \mathbf{a}_n^* &= \Gamma_n^{-1} \gamma_n(h) \implies = \gamma_X(0) + \mathbf{a}_n^{*\top} \Gamma_n \Gamma_n^{-1} \gamma_n(h) - 2\mathbf{a}_n^{*\top} \gamma_n(h) \\ &= \gamma_X(0) - \mathbf{a}_n^{*\top} \gamma_n(h) \end{aligned}$$

Hence, we can conclude that the Mean Squared Predictive Error is obtained by

$$\mathbb{E} \left[(X_{n+h} - P_n X_{n+h})^2 \right] = \gamma_X(0) - \mathbf{a}_n^{*\top} \gamma_n(h)$$

Remark 1.2.

1. For a linear Gaussian process $\{X_t\}$, take $X_{n+h} | X_1, \dots, X_n = P_n X_{n+h} + \epsilon = a_0^* + \mathbf{a}_n^{*\top} \mathbf{x}_n + \epsilon$, then the residual term ϵ must be Gaussian. Then, we have

$$\mathbb{E}[X_{n+h} | X_1, \dots, X_n = P_n X_{n+h}] = \mu - a_0^* + \mu \mathbf{a}_n^{*\top} \mathbf{1} = 0 \xrightarrow{a_0^* = \mu(1 - \mathbf{a}_n^{*\top} \mathbf{1})} \epsilon \sim \mathcal{N}(0, \text{MSPE}(P_n X_{n+h})) = \mathcal{N}(0, \gamma_X(0) - \mathbf{a}_n^{*\top} \gamma_n(h))$$

by our computation; thus, the predictive distribution is also Gaussian:

$$X_{n+h} | X_1, \dots, X_n \sim \mathcal{N}(P_n X_{n+h}, \text{MSPE}(P_n X_{n+h})) = \mathcal{N}\left(a_0^* + \mathbf{a}_n^{*\top} \mathbf{x}_n, \gamma_X(0) - \mathbf{a}_n^{*\top} \gamma_n(h)\right)$$

where $a_0^* = \mu(1 - \mathbf{a}_n^{*\top} \mathbf{1})$ and $\mathbf{a}_n^* = \Gamma_n^{-1} \gamma_n(h)$.

2. For large n , Γ_n^{-1} is obtained through recursive algorithms. As an example look at Durbin-Levinson Algorithm

△

2 Topic: MLE for AR and MA processes

Direct MLE

Suppose that $\{X_t\}$ is a Gaussian time series (process) with mean μ and autocovariance function $\gamma_X(i, j)$. Let $\mathbf{x}_n = (X_1, \dots, X_n) \in \mathbb{R}^n$. Let Γ_n denote the covariance matrix $\Gamma_n = \text{Cov}(\mathbf{x}_n, \mathbf{x}_n)$, and assume that Γ_n is nonsingular. Since any subset of a Gaussian process are jointly Gaussian, the likelihood function for given observation \mathbf{x}_n (a single obseravtion) is

$$L(\Psi) = f(\mathbf{x}_n | \Psi) = \frac{1}{\sqrt{(2\pi)^n |\Gamma_n|}} \exp\left(-\frac{1}{2}(\mathbf{x}_n - \mu)^\top \Gamma_n^{-1}(\mathbf{x}_n - \mu)\right).$$

where Ψ is the parameter set, e.g. for the ARMA model, $\Psi_{ARMA(p,q)} = (\phi_1, \dots, \phi_p, \theta_1, \dots, \theta_q, \sigma^2)$.

To simplify our computation, we could consider the log-likelihood function:

$$\log L(\Psi) = -\frac{n}{2} \log(2\pi) + \frac{1}{2} \log |\Gamma_n|^{-1} - \frac{1}{2}(\mathbf{x}_n - \mu)^\top \Gamma_n^{-1}(\mathbf{x}_n - \mu)$$

Then, the values of parameters Ψ which maximizes the Log-likelihood function is the Maximum Likelihood Estimate of Ψ . That is,

$$\hat{\Psi}_{MLE} = \arg \max_{\Psi} \log L(\Psi)$$

Conditional MLE

For the ease of implementation, we can also use the conditional MLE. Using the fact that

$$f(\mathbf{x}_n | \Psi) = f(x_1, \dots, x_n | \Psi) = f(x_n | x_1, \dots, x_{n-1}, \Psi) f(x_1, \dots, x_{n-1} | \Psi) \stackrel{\dots}{=} f(x_n | x_1, \dots, x_{n-1}, \Psi) f(x_{n-1} | x_1, \dots, x_{n-2}, \Psi) \cdots f(x_1 | \Psi)$$

Then, the MLE is given by

$$\log L(\Psi) = \log f(x_1 | \Psi) + \sum_{i=2}^n \log f(x_i | x_1, \dots, x_{i-1}, \Psi) \quad (2)$$

Similarly, the values of parameters Ψ which maximizes the Log-likelihood function is the Maximum Likelihood Estimate of Ψ . That is,

$$\hat{\Psi}_{MLE} = \arg \max_{\Psi} \log L(\Psi)$$

Remark 2.1.

1. We mostly use Conditional MLE in applications.
2. Since X_1 is the the first data point, the way it's picked is on its own. If X_1 is pick based on some distribution, we can surely keep the term $\log f(x_1 | \Psi)$ in the above equation and work with the exact MLE. If how X_1 is picked is unknown, we can treat it as deterministic, and drop the term $\log f(x_1 | \Psi)$ in the above equation.

△

Problem 2.2. Using a programming language of your choice, repeat the MLE construction and simulation for AR and MA models as discussed this in the class.

Solution:**(a) Conditional MLE for AR(1) Process**

Consider observations $\{X_1, X_2, \dots, X_n\}$ of a stationary Gaussian AR(1) process. For parameter vector is $\Psi = (\phi_0, \phi, \sigma^2)$, we have

$$X_t = \phi_0 + \phi X_{t-1} + Z_t, Z_t \sim IIDN(0, \sigma^2), \quad \mathbb{E}[X_t] = \frac{\phi_0}{1-\phi}, \text{Var}(X_t) = \frac{\sigma^2}{1-\phi^2} \text{ and } \gamma_X(h) = \frac{\phi^h \sigma^2}{1-\phi^2}$$

The probability distribution for X_1 can be written as,

$$X_1 \sim \mathcal{N}\left(\frac{\phi_0}{1-\phi}, \frac{\sigma^2}{1-\phi^2}\right) \Rightarrow f(x_1; \Psi) = \frac{1}{\sqrt{2\pi \frac{\sigma^2}{1-\phi^2}}} \exp\left\{-\frac{\left(x_1 - \frac{\phi_0}{1-\phi}\right)^2}{2 \frac{\sigma^2}{1-\phi^2}}\right\}$$

Since $X_2 = \phi_0 + \phi X_1 + Z_2$, $X_2 | X_1 = x_1 \sim \mathcal{N}(\phi_0 + \phi x_1, \sigma^2)$. The conditional probability distribution for X_2 knowing X_1 can be written as,

$$f(x_2 | x_1; \Psi) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x_2 - \phi_0 - \phi x_1)^2}{2\sigma^2} \right\}$$

Similarly, we have $X_n | X_{n-1} = x_{n-1} \sim \mathcal{N}(\phi_0 + \phi x_{n-1}, \sigma^2)$ and

$$f(x_n | x_{n-1}; \Psi) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x_n - \phi_0 - \phi x_{n-1})^2}{2\sigma^2} \right\}$$

Plug these in Equation (2), we obtain the Log-Likelihood function:

$$\begin{aligned} \log L(\Psi) &= \log f(x_1 | \Psi) + \sum_{i=2}^n \log f(x_i | x_1, \dots, x_{i-1}, \Psi) \\ &= -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log \left(\frac{\sigma^2}{1-\phi^2} \right) - \frac{\left(x_1 - \frac{\phi_0}{1-\phi}\right)^2}{2 \frac{\sigma^2}{1-\phi^2}} + \sum_{i=2}^n \left[-\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma^2) - \frac{(x_i - \phi_0 - \phi x_{i-1})^2}{2\sigma^2} \right] \\ &= -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log \left(\frac{\sigma^2}{1-\phi^2} \right) - \frac{\left(x_1 - \frac{\phi_0}{1-\phi}\right)^2}{2 \frac{\sigma^2}{1-\phi^2}} - \frac{n-1}{2} \log(2\pi) - \frac{n-1}{2} \log(\sigma^2) - \sum_{i=2}^n \frac{(x_i - \phi_0 - \phi x_{i-1})^2}{2\sigma^2} \end{aligned}$$

Maximize this function to estimate the parameters. (Much easier to implement)

Example: Suppose we have AR(1) process $X_t = 0.9X_{t-1} + Z_t$, $Z_t \sim IIDN(0, 0.7^2)$. We simulate this process and pretend that we do not know the true values for ϕ and σ . Then, we use the conditional MLE to estimate the parameter vector $\Psi = (\phi, \sigma^2)$. By the above deduction, we get the log likelihood function is

$$\begin{aligned} \log L(\Psi) &= -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log \left(\frac{\sigma^2}{1-\phi^2} \right) - \frac{x_1^2}{2 \frac{\sigma^2}{1-\phi^2}} - \frac{n-1}{2} \log(2\pi) - \frac{n-1}{2} \log(\sigma^2) - \sum_{i=2}^n \frac{(x_i - \phi x_{i-1})^2}{2\sigma^2} \\ \text{treat } X_1 \text{ as deterministic} &= -\frac{n-1}{2} \log(2\pi) - \frac{n-1}{2} \log(\sigma^2) - \sum_{i=2}^n \frac{(x_i - \phi x_{i-1})^2}{2\sigma^2} \end{aligned}$$

Then, we have

$$\hat{\Psi}_{MLE} = \arg \max_{\Psi} \log L(\Psi)$$

For the coding work on this problem, check the attached jupyter notebook Time Series and Statistical Arbitrage HW2.ipynb

(b) Conditional MLE for MA(1) Process

Consider observations $\{X_1, X_2, \dots, X_n\}$ of a Gaussian MA(1) process. For parameter vector is $\Psi = (\theta_0, \theta_1, \sigma^2)$, we have

$$X_t = \theta_0 + \theta_1 Z_{t-1} + Z_t, \quad Z_t \sim IIDN(0, \sigma^2)$$

The probability distribution for X_t can be written as,

$$X_t | Z_{t-1} \sim \mathcal{N}(\theta_0 + \theta_1 z_{t-1}, \sigma^2) \implies f(x_t | z_{t-1}) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x_t - (\theta_0 + \theta_1 z_{t-1}))^2}{2\sigma^2} \right\}$$

Plug these in Equation (2), we obtain the Log-Likelihood function:

$$\begin{aligned} \log L(\Psi) &= \log f(x_1 | \Psi) + \sum_{i=2}^n \log f(x_i | x_1, \dots, x_{i-1}, \Psi) \\ \text{treat } X_1 \text{ as deterministic} &= -\frac{n-1}{2} \log(2\pi) - \frac{n-1}{2} \log(\sigma^2) - \sum_{i=2}^n \frac{(x_i - \theta_0 - \theta_1 z_{i-1})^2}{2\sigma^2} \end{aligned}$$

Maximize this function to estimate the parameters. (Much easier to implement)

Example: Suppose we have MA(1) process $X_t = 0.5Z_{t-1} + Z_t$, $Z_t \sim IIDN(0, 0.5^2)$. We simulate this process and pretend that we do not know the true values for ϕ and σ . Then, we use the conditional MLE to estimate the parameter vector $\Psi = (\theta_1, \sigma^2)$. By the above deduction, we get the log likelihood function is

$$\log L(\Psi) = -\frac{n-1}{2} \log(2\pi) - \frac{n-1}{2} \log(\sigma^2) - \sum_{i=2}^n \frac{(x_i - \theta_1 z_{i-1})^2}{2\sigma^2}$$

Then, we have

$$\hat{\Psi}_{MLE} = \arg \max_{\Psi} \log L(\Psi)$$

For the coding work on this problem, check the attached jupyter notebook Time Series and Statistical Arbitrage HW2.ipynb

3 Topic: MLE for AR and MA processes

Problem 3.1. For the time series model below:

$$X_t = 0.9X_{t-1} + Z_t \quad Z_t \text{ is } N(0, 0.7^2)$$

- (I) Simulate a path with 1000 data points.
- (II) Take the simulated path as the realized data and estimate the parameters using the MLE estimator.
- (III) Repeat steps (I) and (II) for 100 times and plot the distributions of the parameters you have estimated. What are the 95% confidence levels around your mean estimated parameters?
- (IV) Repeat (III) but use 5000 times instead of 100 times.

Use any programming language or algorithm you are comfortable with but we need to see the code.

Solution: Please check the attached jupyter notebook Time Series and Statistical Arbitrage HW2.ipynb

4 Application: Mean Squared Prediction and the Best Linear Predictor

Problem 4.1. An MA(1) process is given by

$$X_t = Z_t + \theta Z_{t-1} \quad \theta = 0.8, \quad Z_t \sim WN(0, 1)$$

Our observations indicates that $\{X_1 = 3.2020, X_2 = 1.5625\}$.

What is the best linear prediction and MSPE for X_3 ?

Solution: From Equation (1), we know that

$$P_n X_{n+h} := a_0^* + \mathbf{a}_n^{*\top} \mathbf{x}_n \text{ where } \begin{cases} a_0^* &= \mu(1 - \mathbf{a}_n^{*\top} \mathbf{1}) \\ \mathbf{a}_n^* &= \Gamma_n^{-1} \boldsymbol{\gamma}_n(h) \end{cases} \xrightarrow{n=2, h=1} P_2 X_3 := a_0^* + \mathbf{a}_2^{*\top} \mathbf{x}_2 \text{ where } \begin{cases} a_0^* &= \mu(1 - \mathbf{a}_2^{*\top} \mathbf{1}) \\ \mathbf{a}_2^* &= \Gamma_2^{-1} \boldsymbol{\gamma}_2(1) \end{cases}$$

Since $\mathbb{E}[X_t] = 0 = \mu$ and $\gamma_X(h) = (1 + \theta^2)\gamma_Z(h) + \theta\gamma_Z(h-1) + \theta\gamma_Z(h+1) = \begin{cases} (1 + \theta^2)\sigma^2 & \text{if } h = 0 \\ \theta\sigma^2 & \text{if } h = \pm 1 \\ 0 & \text{otherwise} \end{cases}$, we have

$$\Gamma_2 = \begin{bmatrix} \gamma_X(0) & \gamma_X(1) \\ \gamma_X(1) & \gamma_X(0) \end{bmatrix} = \begin{bmatrix} (1 + \theta^2)\sigma^2 & \theta\sigma^2 \\ \theta\sigma^2 & (1 + \theta^2)\sigma^2 \end{bmatrix}, \boldsymbol{\gamma}_2(1) = \begin{bmatrix} \gamma_X(1) \\ \gamma_X(2) \end{bmatrix} = \begin{bmatrix} \theta\sigma^2 \\ 0 \end{bmatrix} \Rightarrow \begin{cases} a_0^* &= 0 \\ \mathbf{a}_2^* &= \Gamma_2^{-1} \boldsymbol{\gamma}_2(1) = \frac{1}{1 + \theta^2 + \theta^4} \begin{bmatrix} \theta + \theta^3 \\ -\theta^2 \end{bmatrix} \xrightarrow{\theta=0.8} \begin{bmatrix} \frac{820}{1281} \\ -\frac{400}{1281} \end{bmatrix} \end{cases}$$

Therefore, the Best Linear Predictor for X_3 is

$$\hat{X}_3 = P_n X_{n+h} := a_0^* + \mathbf{a}_n^{*\top} \mathbf{x}_n = \frac{820}{1281} X_2 - \frac{400}{1281} X_1 = 0.0003512880562062115$$

The corresponding MSPE is

$$MSPE = \mathbb{E} \left[(X_{n+h} - P_n X_{n+h})^2 \right] = \gamma_X(0) - \mathbf{a}_n^{*\top} \boldsymbol{\gamma}_n(h) \xrightarrow{n=2, h=1} \gamma_X(0) - \mathbf{a}_2^{*\top} \boldsymbol{\gamma}_2(1) = 1.64 - \begin{bmatrix} \frac{820}{1281} & -\frac{400}{1281} \end{bmatrix} \begin{bmatrix} 0.8 \\ 0 \end{bmatrix} = 1.1279000780640125$$

Hence, the Best Linear Predictor for X_3 is $\hat{X}_3 = 0.0003512880562062115$ and the corresponding $MSPE = 1.1279000780640125$.